

Christian Hopps
LabN Consulting, LLC

“chopps’ dev” My Linux Kernel Dev Environment

A demo

- Automated Virtual (Munet) Topologies and VMs (qemu)
- Interactive
- Packet Captures
- Debugging with GDB
- Coverage
- Profiling - Flamegraphs
- Tracing

The Dev Source Code (IPTFS)

- `CONFIG_XFRM_IPTFS=y` 😊
- `net/xfrm/xfrm_iptfs.c`
 - non-static functions: `xfrm_iptfs_*`
 - static functions `iptfs_*`
 - 3 sections:
 - State
 - Ingress (output)
 - Egress (input)

The Code

- Linux:
 - Dev Env: <https://github.com/LabNConsulting/iptfs-dev>
 - kernel: <https://github.com/LabNConsulting/iptfs-linux>
 - iproute2: <https://github.com/LabNConsulting/iptfs-iproute2>
 - Dev Env using github actions to run regression/unit tests! 😊
- Testing/Tools:
 - mnet: <https://github.com/LabNConsulting/mnet>
 - wireshark: <https://github.com/LabNConsulting/wireshark>
 - cov.el: <https://github.com/choppsv1/cov/commits/chopps/add-bg-overlay-tint>

Tools

- Hardware/Topology: mUNET
- Testing Framework: pytest
- Used by above
 - Interactive: tmux (best) or SCREEN or X11
 - Profiling: perf and FlameGraph
 - Packet Captures: tshark and Wireshark
 - Debugging: gdb
 - Coverage: native and cov.el

munet – network and hardware emulation

- Config file defines:
 - Topology (node and networks)
 - Kinds (common attributes for nodes)
 - Additional CLI commands
 - yaml/toml/json format
- Launches and configures nodes of the following types:
 - Linux Namespace
 - Container (docker/podman)
 - QemuVM
- Connects nodes together with LAN or P2P networks
 - Normally veth's and bridges
 - Physical Interfaces also supported
 - Mapped into namespaces
 - SRIOV for QemuVM

munet cont.

- Uses TMUX, SCREEN or X11 to open:
 - Extendible Munet CLI
 - Shells in the namespace or containers
 - Consoles to Qemu VMs
 - Tail -f's of stdout/stderr for logs
- Capture packets on any network
- Profile running code (uses perf)
- Converage

```

#                               192.168.0.0/24
#  --+-----+-----+-----+-----+-----+-----+-----+-----+
#    | .1                | .2                | .3                | .4
#  +-----+             +-----+             +-----+             +-----+
#  | h1 | --- net0 --- | r1 | --- net1 --- | r2 | --- net2 --- | r1 |
#  +-----+ .1         .2 +-----+ .2         .3 +-----+ .3         .4 +-----+
#                10.0.0.0/24          10.0.1.0/24          10.0.2.0/24

```

topology:

networks-autonumber: true

networks:

- name: mgmt0
ip: 192.168.0.254/24
- name: net0
ip: 10.0.0.0/24
- name: net1
ip: 10.0.1.0/24
- name: net2
ip: 10.0.2.0/24

nodes:

- name: h1
kind: host
connections: [mgmt0, net0]
- name: r1
kind: linux
connections: [mgmt0, net0, net1]
- name: r2
kind: linux
connections: [mgmt0, net2, net1]
- name: h2
kind: host
connections: [mgmt0, net2]

kinds:

- name: linux
 - merge: ["qemu"]
 - gdb-cmd: "/usr/bin/sudo -E gdb %CONFIGDIR%/../../output-linux/vmlinux"
 - gdb-target-cmds: ["target remote %RUNDIR%/s/gdbserver"]
 - gdb-run-cmds: ["c"]
 - gdb-run-cmd: ["c"]
- qemu:
 - kernel: "%CONFIGDIR%/../../output-linux/arch/x86/boot/bzImage"
 - initrd: "%CONFIGDIR%/../../output-buildroot/images/rootfs.cpio.gz"
 - sshkey: "%CONFIGDIR%/../../root-key"
 - #cmdline-extra: "acpi=off idle=poll nokaslr"
 - #cmdline-extra: "idle=poll nokaslr trace_buf_size=1024M"
 - cmdline-extra: "idle=poll nokaslr"
 - memory: "2048"
 - kvm: true
 - ncpu: 1
 - console:
 - timeout: 180

Setup the Environment

```
$ git clone git@github.com:LabNConsulting/iptfs-dev.git
$ cd iptfs-dev/

$ make setup
[... git clones linux, buildroot, and iproute2]
$ git clone https://github.com/brendangregg/FlameGraph

$ # initialize a python virtual environment
$ python3 -m venv venv
$ . venv/bin/activate
$ pip install -r python-requirements.txt
```

Run interactive

```
$ cd iptfs-dev/  
$ sudo venv/bin/pytest tests/simplenet --pause  
...  
== PAUSING: before step 0: Before test network up ==  
PAUSED, "cli" for CLI, "pdb" to debug, "Enter" to continue: cli  
  
--- Munet CLI Starting ---  
  
munet> con r1  
...
```

Enable coverage and build

```
$ cd iptfs-dev/

$ # change to build coverage kernel
$ sed -i -e 's/^LINUXCONFIG/# &/;/linux-cov.config/s/^# L/L/' Makefile

$ # Enable profiling in linux/net/xfrm
$ echo "GCOV_PROFILE := y" >> linux/net/xfrm/Makefile

$ make
[... builds kernel, buildroot w/ custom iproute]
```

Run test generating coverage

```
$ cd iptfs-dev/  
$ sudo venv/bin/pytest tests/simplenet --coverage  
...  
tests/simplenet/test_simplenet.py::test_policy_tun_up  
-----  
2023-11-01 17:05:34,235 INFO: root: STEP 0: first ping  
2023-11-01 17:05:34,242 INFO: root: STEP 1: second ping  
2023-11-01 17:05:34,247 INFO: root: STEP 2: third ping  
PASSED  
...  
2023-11-01 16:48:15,865 INFO: munet.l3qemuvvm.r1: Saved coverage data in VM at /gcov-data.tgz  
2023-11-01 16:48:15,940 INFO: munet.l3qemuvvm.r1: Saved coverage data on host at /tmp/unet-  
test/tests.simplenet.test_simplenet/r1/gcov-data.tgz  
...
```

Coverage: Extract

```
$ cd iptfs-dev/  
$ scripts/extract-cov.sh  
  
extracting results from /tmp/unet-test/tests.simplenet.test_simplenet/r1/gcov-data.tgz...  
...  
Reading tracefile coverage.info  
Extracting /home/chopps/w/demo/iptfs-dev/linux/net/xfrm/xfrm_iptfs.c  
Extracted 1 files  
Writing data to iptfs.info  
Summary coverage rate:  
  lines.....: 31.7% (350 of 1104 lines)  
  functions...: 43.8% (21 of 48 functions)  
  branches...: no data found  
  
$ # used by emacs cov.el  
$ ln -s $(pwd)/iptfs.info linux/net/xfrm/coverage.info
```

Profiling: Flamegraphs

```
$ cd iptfs-dev/  
$ # enable fast config  
$ sed -i -e 's/^LINUXCONFIG/# &;/linux-fast.config/s/^# L/L/' Makefile  
$ make flame-clean  
$ make flame  
[ runs test, extracts perf data, creates flamegraph and downloads it to my mac ]
```

Debugging

```
$ cd iptfs-dev/  
$ sudo venv/bin/pytest -v tests/simplenet --pause-at-end \  
    --gdb=r1 --gdb-breakpoints=iptfs_enqueue  
...  
[ launches test, opens window with gdb on r1 kernel and sets a breakpoint ]
```


pytest customizations

<code>--profile</code>	Enable profiling if supported by test
<code>--tracing</code>	Enable tracing if supported by test
<code>--cli-on-error</code>	CLI on test failure
<code>--coverage</code>	Enable coverage gathering if supported
<code>--gdb=HOST[,HOST...]</code>	Comma-separated list of nodes to launch gdb on, or 'all'
<code>--gdb-breakpoints=BREAKPOINT[,BREAKPOINT...]</code>	Comma-separated list of breakpoints
<code>--gdb-use-emacs</code>	Use emacsclient to run gdb instead of a shell
<code>--pcap=NET[,NET...]</code>	Comma-separated list of networks to capture packets on, or 'all'
<code>--pause</code>	Pause after each test
<code>--pause-at-end</code>	Pause before taking munet down
<code>--shell=NODE[,NODE...]</code>	Comma-separated list of nodes to spawn shell on, or 'all'
<code>--stdout=NODE[,NODE...]</code>	Comma-separated list of nodes to open tail-f stdout window on, or 'all'
<code>--stderr=NODE[,NODE...]</code>	Comma-separated list of nodes to open tail-f stderr window on, or 'all'

Questions and Comments
